

Data Preprocessing and Collection:

We ran a simple line of code to check the number of missing values in the dataset. On running the line of code, we found that the value it returned was 'False' indicating that there were no missing values in the dataset.

```
In [10]: # Checking if there are any null values in the dataframe  
data.isnull().any().any()
```

```
Out[10]: False
```

```
In [11]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 7043 entries, 0 to 7042  
Data columns (total 21 columns):  
#   Column                Non-Null Count  Dtype  
---  ----  
0   customerID            7043 non-null   object  
1   gender                 7043 non-null   object  
2   SeniorCitizen          7043 non-null   int64  
3   Partner                7043 non-null   object  
4   Dependents             7043 non-null   object  
5   tenure                 7043 non-null   int64  
6   PhoneService           7043 non-null   object  
7   MultipleLines           7043 non-null   object  
8   InternetService        7043 non-null   object  
9   OnlineSecurity         7043 non-null   object  
10  OnlineBackup           7043 non-null   object  
11  DeviceProtection       7043 non-null   object  
12  TechSupport            7043 non-null   object  
13  StreamingTV            7043 non-null   object
```

```
In [12]: data.shape
```

```
Out[12]: (7043, 21)
```

In the screenshot below you can see that we are dropping the column 'CustomerID' because it is of no relevance in our study. Along with that we also run a line of python code to filter the data frame to only include rows where the value in the column 'TotalCharges' is an empty string.

```
In [14]: # Dropping the following column because it serves no purpose in our analysis #
data = data.drop(["customerID"], axis = 1)
data.head()
```

```
Out[14]:
```

| | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | OnlineBackup |
|---|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|--------------|
| 0 | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | Yes |
| 1 | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | No |
| 2 | Male | 0 | No | No | 2 | Yes | No | DSL | Yes | Yes |
| 3 | Male | 0 | No | No | 45 | No | No phone service | DSL | Yes | No |
| 4 | Female | 0 | No | No | 2 | Yes | No | Fiber optic | No | No |

```
In [15]: data[data["TotalCharges"] == ' ']
```

```
Out[15]:
```

| | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | OnlineBackup |
|------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|---------------------|---------------------|
| 488 | Female | 0 | Yes | Yes | 0 | No | No phone service | DSL | Yes | No |
| 753 | Male | 0 | No | Yes | 0 | Yes | No | No | No internet service | No internet service |
| 936 | Female | 0 | Yes | Yes | 0 | Yes | No | DSL | Yes | No |
| 1082 | Male | 0 | Yes | Yes | 0 | Yes | Yes | No | No internet service | No internet service |

In this screenshot, we run a code to convert the column 'TotalCharges' to a numeric datatype using the Pandas "to numeric" function to perform certain numerical calculations on the column. The argument "errors = coerce" is asking the function to convert any non-numeric values in the column to "NaN" (Not a Number) values. You will notice that on running this code, we understand that there are 11 missing records in the column 'TotalCharges'.

```
In [16]: data['TotalCharges'] = pd.to_numeric(data.TotalCharges, errors='coerce')
data.isnull().sum()
### you will notice 11 missing records in the 'totalcharges' column
```

```
Out[16]:
```

| | |
|------------------|----|
| gender | 0 |
| SeniorCitizen | 0 |
| Partner | 0 |
| Dependents | 0 |
| tenure | 0 |
| PhoneService | 0 |
| MultipleLines | 0 |
| InternetService | 0 |
| OnlineSecurity | 0 |
| OnlineBackup | 0 |
| DeviceProtection | 0 |
| TechSupport | 0 |
| StreamingTV | 0 |
| StreamingMovies | 0 |
| Contract | 0 |
| PaperlessBilling | 0 |
| PaymentMethod | 0 |
| MonthlyCharges | 0 |
| TotalCharges | 11 |
| Churn | 0 |
| dtype: int64 | |

In this screenshot, we see that we are carrying out the process of imputation. This helps with rectifying or filling the missing values or missing records in the column "TotalCharges". The purpose of this code is to fill any missing record in the column with a reasonable estimate based on the mean value of the column. This helps with ensuring that the dataset is complete and can be used for further analysis.

```

In [18]: ### to rectify the missing values in the dataframe we are going to impute the missing records with the mean value of t
mean_value = data['TotalCharges'].mean()
data['TotalCharges'].fillna(mean_value, inplace=True)
data.head()

Out[18]:

```

| | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | OnlineBackup | DeviceProtection | TechSupport |
|---|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|--------------|------------------|-------------|
| 0 | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | Yes | No | N |
| 1 | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | No | Yes | N |
| 2 | Male | 0 | No | No | 2 | Yes | No | DSL | Yes | Yes | No | N |
| 3 | Male | 0 | No | No | 45 | No | No phone service | DSL | Yes | No | Yes | Ye |
| 4 | Female | 0 | No | No | 2 | Yes | No | Fiber optic | No | No | No | N |

```

In [19]: ## You can see that the missing values have been imputed.
data.isnull().sum()

Out[19]:
gender                0
SeniorCitizen         0
Partner               0
Dependents            0
tenure                0
PhoneService          0
MultipleLines         0
InternetService       0
OnlineSecurity        0
OnlineBackup          0
DeviceProtection     0
TechSupport           0
StreamingTV           0
StreamingMovies       0
Contract              0
PaperlessBilling      0
PaymentMethod         0
MonthlyCharges        0
TotalCharges          0
Churn                 0
dtype: int64

```

In this, we notice a correlation matrix that is plotted as a heatmap indicating the correlation between the different features in the dataset. With this plot we can explore the different features and come up with various visualization plots to illustrate the key inferences we can draw from the dataset.


Plotting a heatmap to show the correlation between all the columns in the dataframe

```

In [21]: import seaborn as sns

corr = data.corr()
sns.heatmap(corr, annot=True)
py.show()

```



The lines of code below are executed to perform feature scaling on the columns specified. Feature scaling is a very common data preprocessing technique that is used to normalize the range of feature values in a dataset. This is usually done to ensure that

the features with larger ranges of values do not dominate features with smaller ranges during model training. This is the last preprocessing step that is executed before attaching all the data mining models to train the dataset which will be explained in the next section.